

UMAR AHMED THAMEEM AHMED

umarahmed.thameemahmed@colorado.edu | (720) 232-0077 | [linkedin.com/in/umarahmed2000](https://www.linkedin.com/in/umarahmed2000) | github.com/UmarAhmed30 | umarahmed.codes

EDUCATION

University of Colorado Boulder

Boulder, CO

Master of Science – Computer Science

Aug 2025 – May 2027

- Coursework: Distributed Systems, Systems for ML, Datacenter Scale Computing, Big Data Architecture
- CGPA: 4.0/4.0

College of Engineering Guindy, Anna University

Chennai, India

Bachelor of Engineering – Computer Science and Engineering

Aug 2018 – Apr 2022

- Coursework: Data Structures & Algorithms, Database Systems, Operating Systems, Networks, Object-Oriented Programming
- CGPA: 8.75/10 (First Class with Distinction)

WORK EXPERIENCE

F5, Inc.

Boulder, CO

Software Development Engineer Intern

May 2026 – Aug 2026

Tech Stack: Go, gRPC, Kubernetes, Terraform, Bash, AWS

- Architect end-to-end, cost-efficient AWS deployments for NGINX SaaS by delivering autoscaling support with AWS Auto Scaling Groups. Build usage-based billing tracking, observability dashboards, and custom UI configurations

Nueromind Technologies

Bangalore, India

Founding Software Engineer

Jan 2024 – Jul 2025

Tech Stack: Python, React, MongoDB, Docker, FFmpeg, AWS

- Reduced client onboarding time for CCTV integrations by **40%** by building a scalable multi-protocol video ingestion pipeline supporting WebRTC, RTSP and HLS. Strengthened SDLC practices through code reviews and technical documentation
- Optimized the frame processing pipeline by integrating HLS-based streaming to deliver processed video output **3x** faster to users
- Slashed video preprocessing latency from **180s** to **50s** across **~150GB/day** of video data by revamping the FFmpeg pipeline into a maintainable, multithreaded, I/O-tuned architecture with asynchronous S3 prefetching for real-time WebRTC streams
- Improved object detection accuracy by **20%** by building a human-in-the-loop annotation system with React and FastAPI for reviewing and correcting model predictions

Infibeam Avenues

Bangalore, India

Software Development Engineer

Jul 2022 – Jan 2024

Tech Stack: Ruby on Rails, AngularJS, MySQL, Docker, ELK, Kubernetes

- Shipped a Firebase-based seller notification system serving **1M+** sellers, driving a **15%** increase in seller engagement through real-time push notifications within a cross-functional Agile team
- Saved support teams **120+ hours/month** by spearheading a suite of REST APIs in Ruby on Rails to automate seller ticketing workflows, collaborating with third-party partners and integrating external services
- Enabled admins to process **2K+** seller stock approval requests per day by delivering a pull-based work allocation system. Distributed requests across admins and documented allocation logic, request states, and edge cases for support teams

Totallr Technologies

Chennai, India

Software Development Intern

Apr 2021 – Jun 2021

Tech Stack: Angular, Spring Boot, Java, MySQL, AWS

- Shipped a full-stack Point-of-Sale system using Spring Boot and Angular, reducing checkout times and automating invoicing across **25+** enterprise clients at launch
- Identified and optimized slow SQL queries, cutting latency from **1100 ms** to **600 ms** under peak transaction loads through refactored access patterns, triggers, stored procedures, and rigorous testing

PROJECTS

Distributed Log Processing System | Next.js, FastAPI, ClickHouse, GCP

[\[Code\]](#)

- Accelerated end-to-end log analysis throughput by **10x** to **~1M events/min** compared to a synchronous baseline by engineering an Apache Kafka + Flink ETL pipeline and an observability dashboard with <100ms query latency via Redis caching

HyRA - Hybrid Routing Agent | React, Flask, MySQL, GCP

[\[Code\]](#)

- Developed a cost-aware LLM routing service with vLLM and LangGraph, improving routing accuracy from **71%** to **89%** on a curated evaluation set through an LLM-as-a-Judge feedback loop and iterative prompt tuning
- Productionized model serving on GCP with Langfuse-based monitoring and performance profiling for distributed LLM workloads

SKILLS

Languages & Databases: Python, Ruby, Java, TypeScript/JavaScript, MySQL, PostgreSQL, MongoDB, Redis

Web & App Development: HTML/CSS, React, Angular, Ruby on Rails, Flask, FastAPI

Cloud & DevOps: AWS, GCP, Docker, Kubernetes, CI/CD, Terraform, Nginx, ELK Stack, Prometheus, Grafana, Linux

LLM & Developer Tools: OpenAI API, Claude, Copilot, Visual Studio, Cursor, Git/GitHub, Postman